

# Pass-Through, Welfare, and Incidence under Product Returns in Perfect Competition

Draft notes

July 24, 2026

## 1 Introduction

About one in five products sold online is returned. Pass-through, tax incidence, and the efficiency cost of taxation are well understood when purchase settles who keeps the good. In that classical setting, the demand and supply curves are sufficient for consumer surplus, producer surplus, and deadweight loss. With returns, allocation is decided in two stages. Buyers first choose whether to purchase, then whether to keep or return after receiving the item. Never buying and buying then returning leave the same final allocation—the consumer does not keep the good—yet their welfare costs differ. A returned purchase incurs shipping, hassle, and handling that a non-purchase avoids, so demand and supply alone omit an economically relevant margin.

What objects then replace the classical demand and supply curves, and how misleading is an analysis that ignores returns? We study these questions under perfect competition with product returns and seller handling costs. We characterize pass-through, incidence, and welfare losses with returns, compare them to the no-return benchmark, and identify where omitting the return margin changes conclusions.

Consumer surplus depends only on two aggregates, gross demand and the return rate. Heterogeneity in tastes, match values, and return costs enters welfare solely through these aggregates. Producer surplus retains the usual geometry—area to the left of the supply curve—but is evaluated at net revenue, list price minus expected refunds and handling, rather than at list price.

Because consumers pay the list price while firms respond to net revenue, pass-through from cost shocks into list prices is governed by how net revenue moves with the list price. Tax neutrality fails for a different reason: the keep-or-return decision depends on the refund price, which differs according to which party remits the tax. An analyst who ignores these two differences and applies the classical formulas therefore misstates pass-through, incidence, and deadweight loss whenever the return margin is active. We show that pass-through of a one-dollar production-cost shock equals that of a seller-remitted unit tax, but differs under buyer remittance. The ratio of handling to production-cost pass-through is governed only by the aggregate return rate. We also give an easy check, from the return-rate elasticity and the ratio of handling to list price, for when the classical no-return pass-through formula over- or underestimates the correct returns-adjusted pass-through.

## 2 Setting

We begin with the market primitives: purchase and return behavior on the consumer side, and competitive supply with return handling on the firm side. A homogeneous product is sold at list price  $p$ . Consumers indexed by  $i$  differ in return hassle  $\kappa_i$ , the disutility from returning the product,

and in match distribution  $G_i$  over post-purchase match quality  $u$ . After purchase, type  $i$  draws  $u \sim G_i$  and keeps the item whenever  $u - p \geq -\kappa_i$ , earning surplus  $u - p$ , and returns it otherwise, earning  $-\kappa_i$ .

Buying at price  $p$  yields expected surplus

$$s_i(p) = \mathbb{E}_{u \sim G_i} [\max(u - p, -\kappa_i)].$$

On the purchase margin, all consumer-side heterogeneity is summarized by willingness to pay  $v_i$ , defined by  $s_i(v_i) = 0$ , so type  $i$  buys if and only if  $p \leq v_i$ . Across types, willingness to pay has distribution  $F(v) = \Pr(v_i \leq v)$  with density  $f$ , so gross demand is  $D(p) = 1 - F(p)$ . Of the units sold at price  $p$ , a share  $r(p)$  are returned.

To understand how  $r(p)$  changes with price, write type  $i$ 's return probability as  $\rho_i(p) = \Pr(u < p - \kappa_i)$ . Willingness to pay is sufficient for purchase but not for returns: types with the same  $v_i$  may still differ in  $(\kappa_i, G_i)$ . Averaging within WTP cells,

$$\rho(v, p) \equiv \mathbb{E}[\rho_i(p) \mid v_i = v],$$

so the aggregate return rate is the truncated mean

$$r(p) = \mathbb{E}[\rho(v_i, p) \mid v_i \geq p].$$

Differentiating yields

$$r'(p) = \mathbb{E}[\partial \rho(v_i, p) / \partial p \mid v_i \geq p] + \frac{f(p)}{D(p)} (r(p) - \rho(p, p)),$$

where  $\rho(p, p)$  is the average return probability among types at the purchase margin  $v_i = p$ .<sup>1</sup> The first term is behavioral: a higher keep threshold raises return probabilities among those who still buy. The second is selection: a higher price drops cutoff types at the purchase margin, and that term would be negative if those cutoff types return more than the average purchaser.<sup>2</sup> Behavioral and selection effects can therefore pull in opposite directions, and we do not impose a sign on  $r'(p)$ .

The product market is perfectly competitive. Producer  $j$  takes  $p$  as given, chooses output at cost  $C_j(q_j)$ , and incurs a net return handling cost  $h$  on each returned unit, equal to gross handling cost minus the scrap value of the returned product. This net cost may be negative, but we maintain  $p + h > 0$ , which holds whenever scrap value does not exceed list price.

### 3 Consumer side

We turn first to consumer surplus and show that only the aggregates  $D$  and  $r$  enter. Write  $s(v, p)$  for expected surplus at list price  $p$  among types with willingness to pay  $v$ , so  $s(v, v) = 0$ . Consumer surplus integrates surplus over all types who buy,

$$CS(p) = \int_p^\infty s(v, p) dF(v). \tag{1}$$

---

<sup>1</sup>With  $\rho(v, p)$  as defined in the text,  $r(p) = D(p)^{-1} \int_p^\infty \rho(v, p) dF(v)$ . Differentiating and using  $D'(p) = -f(p)$  gives

$$r'(p) = D(p)^{-1} \int_p^\infty \partial \rho(v, p) / \partial p dF(v) + \frac{f(p)}{D(p)} (r(p) - \rho(p, p)),$$

which is the display in the text.

<sup>2</sup>Selection need not be negative for arbitrary  $(\kappa_i, G_i)$ . It is signed under an additive match specification  $u = v_i + \varepsilon$  with shock  $\varepsilon$  of any distribution independent of willingness to pay, and with return hassle likewise independent of  $v$ : then  $\rho(v, p)$  decreases in  $v$ , so  $\rho(p, p) \geq r(p)$  and the selection term is non-positive.

Differentiating in  $p$  gives an extensive-margin term from types entering or leaving the purchase margin and an intensive-margin term from lower surplus among inframarginal buyers,

$$CS'(p) = -s(p, p) f(p) + \int_p^\infty s_p(v, p) dF(v),$$

where  $s_p(v, p) \equiv \partial s(v, p)/\partial p$ . On the extensive margin,  $s(p, p) = 0$ , so entry and exit contribute nothing at first order. On the intensive margin,  $s_p(v, p)$  equals minus the keep probability, because a higher price harms the consumer only when she keeps the purchase,

$$s_p(v, p) = -\Pr(u \geq p - \kappa_i).$$

Aggregating over buyers who purchase then gives

$$CS'(p) = \int_p^\infty s_p(v, p) dF(v) = -D(p)(1 - r(p)),$$

by the definitions of  $D(p)$  and  $r(p)$ . A small price change therefore satisfies

$$dCS = -D(p)(1 - r(p)) dp. \quad (2)$$

The extensive- and intensive-margin decomposition makes the local formula transparent. Consumer surplus changes depend only on the aggregates  $D(p)$  and  $r(p)$ , both at each price and integrated along any path,

$$\Delta CS = - \int_{p_0}^{p_1} D(p)(1 - r(p)) dp.$$

Micro heterogeneity in price responses can be rich. Price  $p$  enters each  $s_i(p)$  nonlinearly and can shift participation, return decisions, and inframarginal surplus differently across types. Yet once  $(D, r)$  are given, this microstructure is irrelevant for  $dCS$  and  $\Delta CS$ . Two economies with the same paths of  $D(p)$  and  $r(p)$  deliver the same consumer surplus loss, whether high- $v_i$  types adjust return behavior more steeply than low- $v_i$  types or not. Only kept units enter at first order, through  $1 - r(p)$ , not gross purchases  $D(p)$  alone.

Let  $\bar{p}$  denote the highest willingness to pay in the population, so  $D(\bar{p}) = 0$  and  $CS(\bar{p}) = 0$ . Integrating the local formula then rewrites consumer surplus as area under kept volume,

$$CS(p) = \int_p^{\bar{p}} D(z)(1 - r(z)) dz. \quad (3)$$

## 4 Supply side

On the firm side, returns make list price an imperfect summary of producer payoffs. A gross sale brings revenue  $p$  if the consumer keeps the item and costs refund  $p$  plus net handling  $h$  if the item is returned. Among gross sales at list price  $p$ , share  $r(p)$  are returned, so expected revenue per gross unit—the *effective marginal revenue*—is

$$\mu(p, h) \equiv p - r(p)(p + h). \quad (4)$$

Under perfect competition, firm  $j$  takes  $p$  as given and therefore faces a fixed number  $\mu$  in its supply decision. At equilibrium,  $\mu = \mu(p, h)$ .

Given  $\mu$ , firm  $j$  chooses output to solve

$$\max_{q_j \geq 0} \mu q_j - C_j(q_j), \quad \text{FOC: } C'_j(q_j) = \mu \text{ (if } q_j > 0),$$

with  $C_j(0) = 0$ . Write  $q_j(\mu)$  for firm  $j$ 's supply, set to zero when  $\mu$  is below its minimum marginal cost. We assume each  $C_j$  has strictly increasing marginal cost, so  $q'_j(\mu) > 0$  for active firms.

Summing across firms defines market supply and industry profit,

$$S(\mu) \equiv \sum_j q_j(\mu), \quad PS(\mu) \equiv \sum_j [\mu q_j(\mu) - C_j(q_j(\mu))]. \quad (5)$$

Under that assumption,  $S'(\mu) > 0$ . Market supply therefore rises with effective revenue, but not necessarily with list price. Firms care about  $\mu$ , and  $\mu(p, h)$  can fail to rise with  $p$  when refunds and handling eat into the list-price gain. Which firms are active at each  $\mu$  determines both the market quantity and total profit.

Differentiating  $PS(\mu)$  with respect to  $\mu$  gives, for each firm,

$$\frac{d}{d\mu} [\mu q_j - C_j(q_j)] = q_j + q'_j(\mu) [\mu - C'_j(q_j)].$$

Active firms satisfy  $C'_j(q_j) = \mu$  at their chosen output, so the adjustment term  $q'_j(\mu) [\mu - C'_j(q_j)]$  is zero for each firm—equivalently, the envelope theorem applied to each firm's profit problem gives  $d\pi_j/d\mu = q_j(\mu)$ . Summing over firms,

$$PS'(\mu) = \sum_j q_j(\mu) = S(\mu).$$

With no production at  $\mu = 0$ , it follows that industry profit equals the area to the left of the market supply curve,

$$PS(\mu) = \int_0^\mu S(z) dz. \quad (6)$$

At equilibrium  $\mu = \mu(p, h)$ ,  $PS = \int_0^{\mu(p, h)} S(z) dz$ .

## 5 Failure of tax neutrality

With firms responding to  $\mu$  rather than list price, a natural question is whether remittance of a unit tax remains neutral. Weyl and Fabinger (2013) show that without returns, buyer and seller remittance deliver the same equilibrium. With returns, we ask whether that equivalence survives by writing the two clearing conditions at a common checkout price  $p$ —the total payment at the register—and a per-unit tax  $\tau$ . Fix  $p$  and  $\tau$ . Seller remittance clears at

$$D(p) = S(p - \tau - r(p)(p + h)), \quad (7)$$

because both the return rate and per-return cost  $p + h$  are evaluated at checkout price  $p$ . Buyer remittance clears at

$$D(p) = S((p - \tau) - r(p - \tau)(p - \tau + h)), \quad (8)$$

because the seller posts list price  $p - \tau$ , so returns and refund cost are evaluated there. The two conditions define the same equilibrium only if the arguments of  $S(\cdot)$  coincide. Canceling  $p - \tau$  yields

$$r(p)(p + h) = r(p - \tau)(p - \tau + h). \quad (9)$$

**Result 1** (Failure of tax neutrality with returns). *Seller and buyer taxes define the same equilibrium if and only if  $r(p)(p + h) = r(p - \tau)(p - \tau + h)$ .*<sup>3</sup>

Tax neutrality breaks down because the refunded amount—and therefore the keep-or-return decision—depends on which party remits the tax. Under seller remittance the tax-inclusive checkout price is refunded; under buyer remittance only the pre-tax list price is. Purchase is always based on checkout price  $p$ , but returns are evaluated at  $p$  under seller remittance and at  $p - \tau$  under buyer remittance, so effective revenue generally differs and the two clearing conditions are not equivalent.

## 6 Pass-through

Equilibrium satisfies  $D(p) = S(\mu(p, h))$ . Write  $Q \equiv D(p)$ . We characterize pass-through into list price for three shocks: a uniform one-dollar parallel shift in every firm's marginal production cost (indexed by  $c$ ), a one-dollar increase in handling cost  $h$ , and a unit tax under seller and buyer remittance. Production cost shifts market supply at a given effective margin  $\mu$ , while handling lowers  $\mu(p, h)$  at a fixed list price  $p$ .

For production and handling, write the clearing condition as

$$m(p, x) \equiv D(p) - S(\mu(p, x), x) = 0, \quad x \in \{c, h\}. \quad (10)$$

Differentiating while keeping  $m = 0$  defines pass-through  $\rho_x \equiv dp/dx$  by

$$\rho_x = -\frac{m_x}{m_p}, \quad (11)$$

and expanding  $m = D - S$  gives

$$m_p = D'(p) - S'(\mu) \mu_p, \quad m_x = -S'(\mu) \mu_x - S_x, \quad (12)$$

hence

$$\rho_x = \frac{-S'(\mu) \mu_x - S_x}{S'(\mu) \mu_p - D'(p)}. \quad (13)$$

The two shocks differ in which margin moves at fixed  $p$  or fixed  $\mu$ ,

$$\mu_c = 0, \quad S_c = -S'(\mu), \quad \mu_h = -r(p), \quad S_h = 0. \quad (14)$$

The production term is  $S_c = -S'(\mu)$  because a one-dollar parallel rise in every firm's marginal cost is equivalent, at fixed  $\mu$ , to evaluating market supply at  $\mu - 1$ . Handling moves  $\mu$  by only  $r(p)$  per dollar relative to that unit production shift, so its numerator in (13) is the fraction  $r(p)$  of the production numerator.

Turning to the unit tax, failure of tax neutrality requires separate checkout prices  $p^s$  and  $p^b$  under seller and buyer remittance. Write  $\rho_\tau^s \equiv dp^s/d\tau$  and  $\rho_\tau^b \equiv dp^b/d\tau$ . At  $\tau = 0$ , (13) with  $S_\tau = 0$  gives

$$\rho_\tau^s = \frac{-S'(\mu) \mu_\tau^s}{S'(\mu) \mu_p - D'(p^s)}, \quad \rho_\tau^b = \frac{-S'(\mu) \mu_\tau^b}{S'(\mu) \mu_p - D'(p^b)}. \quad (15)$$

The regimes differ only through  $\mu_\tau^s$  and  $\mu_\tau^b$  at fixed checkout price. At  $\tau = 0$ ,  $\mu_\tau^s = -1$ . Under buyer remittance,  $\tau$  enters effective revenue in the same way as  $-p$ , so  $\mu_\tau^b = -\mu_p$ . Seller remittance therefore has the same format as a production-cost shock, while buyer remittance scales that format by  $\mu_p$ . Collecting these cases, we have

---

<sup>3</sup>Neutrality requires the return wedge to be the same at checkout price  $p$  and at posted list price  $p - \tau$ . That holds for all taxes only if  $r(p)(p + h)$  is constant in posted list price, which requires  $r(p) = C/(p + h)$  for some constant  $C$ ; without that functional form on  $r$ , the two clearing conditions need not define the same equilibrium.

**Result 2** (Pass-through into list price). *Production-cost pass-through is*

$$\rho_c = \frac{S'(\mu)}{S'(\mu)\mu_p - D'(p)}. \quad (16)$$

The remaining shocks connect to this object by

$$\rho_h = r(p)\rho_c, \quad \rho_\tau^s = \rho_c, \quad \rho_\tau^b = \mu_p\rho_c. \quad (17)$$

When  $r \equiv 0$ , one has  $\mu_p = 1$  and therefore  $\rho_c = \rho_\tau^s = \rho_\tau^b$ , which equals the textbook competitive pass-through.

## 6.1 Elasticity representation

The same pass-through formulas can be rewritten in elasticities, which makes the role of  $\mu_p$  transparent. Write  $\mu_p \equiv d\mu/dp$  for the response of effective revenue to list price, and define demand and supply elasticities with respect to list price as usual by

$$\varepsilon_D(p) \equiv \frac{pD'(p)}{D(p)} < 0, \quad \varepsilon_S(p) \equiv \frac{dS(\mu(p))}{dp} \frac{p}{S(\mu(p))} = S'(\mu)\mu_p \frac{p}{S(\mu)}.$$

Here  $\varepsilon_S(p)$  is the observed elasticity of market supply with respect to the clearing price. Under  $S'(\mu) > 0$ , it has the sign of  $\mu_p$ .

Without returns, supply is a function of list price, so  $S' = dS/dp$ . A one-dollar cost increase cuts supply by  $S'$  units at the original price, creating an excess-demand gap of that size. Raising list price by one dollar closes the gap from both sides. Demand falls by  $|D'|$  and supply recovers by  $S'$ . Closing a gap of size  $S'$  therefore requires a price rise of  $S'/(S' + |D'|)$ , so standard pass-through equals  $\varepsilon_S/(-\varepsilon_D + \varepsilon_S)$ .

With returns, supply is a function of effective revenue, so  $S'(\mu) = dS/d\mu$  rather than  $dS/dp$ . A one-dollar production-cost increase is equivalent to lowering  $\mu$  by one dollar, so at fixed list price supply falls by  $S'(\mu)$  units, again creating an excess-demand gap of that size. Raising list price by one dollar closes the gap at rate  $S'(\mu)\mu_p - D'(p)$ , so

$$\rho_c = \frac{S'(\mu)}{S'(\mu)\mu_p - D'(p)}.$$

The cost shock lowers  $\mu$  by one dollar, while a dollar of list price raises  $\mu$  by  $\mu_p$ . A higher  $\mu_p$  means a one-dollar price increase restores more effective revenue—and therefore more quantity supplied—than the one-dollar cost shock removes, so a smaller list-price movement closes the excess-demand gap. Measuring supply against list price through  $\varepsilon_S$  puts the factor  $1/\mu_p$  in the numerator,

$$\rho_c = \frac{\varepsilon_S(p)/\mu_p}{-\varepsilon_D(p) + \varepsilon_S(p)}. \quad (18)$$

Pass-through is therefore dampened relative to the no-return benchmark when  $\mu_p > 1$  and amplified when  $\mu_p < 1$ . If returns are absent,  $r \equiv 0$  forces  $\mu_p = 1$ , and (18) collapses to the standard formula.

Whether  $\mu_p$  exceeds one depends on how steeply the return rate falls with price. Differentiating  $\mu = p - r(p)(p + h)$  gives

$$\mu_p = 1 - r(p) - r'(p)(p + h).$$

Write  $\varepsilon_r(p) \equiv p r'(p)/r(p)$  for the elasticity of the return rate in list price. Then  $\mu_p > 1$  if and only if  $\varepsilon_r(p) < -p/(p + h)$ . That is, selection must outweigh the behavioral rise in individual return probabilities by enough that aggregate  $r$  falls steeply in  $p$ . When instead the behavioral response dominates,  $\varepsilon_r$  is less negative (or positive) and  $\mu_p < 1$ . Comparing  $\mu_p$  with the classical benchmark  $\mu_p = 1$  then signs the bias for an analyst who ignores returns.

**Result 3.** *Relative to the correct returns-adjusted  $\rho_c$ , the classical no-return pass-through formula overestimates pass-through when  $\varepsilon_r(p) < -p/(p+h)$  and underestimates it otherwise.*

The practical value is that the bias can be signed without estimating the full returns-adjusted model. The return-rate elasticity and the ratio of handling to list price are enough. The economics behind the threshold is familiar from the selection and behavioral forces already in  $\varepsilon_r$ . When selection dominates, a higher list price lowers the return rate, so the refund-and-handling drag on effective revenue shrinks as price rises. Effective revenue can then rise by more than one dollar per dollar of list price, and a smaller price movement closes the excess-demand gap than the classical formula predicts. Costly handling strengthens that channel and makes overestimation more likely. Otherwise, effective revenue rises by less than one dollar per dollar of list price, so list price must move by more than the classical formula predicts to close the gap, and the classical no-return formula underestimates pass-through.

## 7 Incidence

Pass-through into list price determines the welfare loss from each cost shifter and how that loss is split between consumers and producers. Throughout,  $Q \equiv D(p)$  and  $r \equiv r(p)$  denote the equilibrium gross quantity and the return rate. From the consumer-side analysis, a small increase in any shifter  $x$  changes consumer surplus by

$$\frac{dCS}{dx} = -Q(1-r)\rho_x, \quad (19)$$

where  $\rho_x \equiv dp/dx$  is the pass-through of the shock into the list price.

Producer surplus is less immediate because a shock may shift the supply schedule directly as well as indirectly through the effective marginal revenue  $\mu$ . Let  $S(\mu, x)$  be the market supply curve, where the second argument allows for a direct effect of  $x$  on firms' costs, and write producer surplus as  $PS(\mu, x) = \int_0^\mu S(z, x) dz$ . Differentiating with respect to  $x$ ,

$$\frac{dPS}{dx} = S(\mu, x) \frac{d\mu}{dx} + \int_0^\mu \frac{\partial S(z, x)}{\partial x} dz.$$

Using the equilibrium condition  $S(\mu, x) = Q$  and denoting  $\Sigma_x \equiv \int_0^\mu (\partial S(z, x)/\partial x) dz$ , we obtain

$$\frac{dPS}{dx} = Q \frac{d\mu}{dx} + \Sigma_x = Q(\mu_p \rho_x + \mu_x) + \Sigma_x, \quad (20)$$

with  $\mu_p \equiv \partial\mu/\partial p$  and  $\mu_x \equiv \partial\mu/\partial x$  evaluated at fixed list price.

The term  $\Sigma_x$  captures the direct shift of the supply curve holding  $\mu$  constant. A uniform, parallel increase in every firm's marginal production cost by one dollar ( $x = c$ ) moves the supply curve inward according to  $\partial S(\mu, c)/\partial c = -S'(\mu)$ , so  $\Sigma_c = \int_0^\mu (-S'(z)) dz = -S(\mu) = -Q$ . For a handling-cost shock  $h$  or a per-unit tax (whether seller- or buyer-remitted), the supply schedule at a given  $\mu$  is unchanged, hence  $\Sigma_h = \Sigma_{\tau_s} = \Sigma_{\tau_b} = 0$ . The partial derivatives  $\mu_x$  are obtained directly from the definition of effective marginal revenue:  $\mu_c = 0$ ,  $\mu_h = -r$ ,  $\mu_{\tau_s} = -1$ , and, for the buyer-remitted tax evaluated at  $\tau = 0$ ,  $\mu_{\tau_b} = r + (p+h)r' - 1$  as derived in Section 5.

Substituting these values together with the pass-through expressions  $\rho_h = r\rho_c$  and  $\rho_\tau^s = \rho_c$  into

(20) yields the producer-surplus responses:

$$\frac{dPS}{dc} = Q(\mu_p \rho_c - 1), \quad (21)$$

$$\frac{dPS}{dh} = Q(\mu_p \rho_h - r), \quad (22)$$

$$\frac{dPS}{d\tau_s} = Q(\mu_p \rho_c - 1), \quad (23)$$

$$\frac{dPS}{d\tau_b} = Q(\mu_p \rho_\tau^b + \mu_{\tau_b}). \quad (24)$$

The incidence ratio  $I(x) \equiv |dCS/dx| / |dPS/dx|$  measures how the total loss is divided. For production cost, handling cost, and the seller-remitted tax, the ratio simplifies to the same expression:

$$I(c) = I(h) = I(\tau_s) = \frac{(1-r)\rho_c}{1 - \mu_p \rho_c}. \quad (25)$$

For the buyer-remitted tax, the ratio is

$$I(\tau_b) = \frac{(1-r)\rho_\tau^b}{-(\mu_p \rho_\tau^b + \mu_{\tau_b})}, \quad \mu_{\tau_b} = r + (p+h)r' - 1. \quad (26)$$

Equation (25) generalises the standard competitive-market formula  $I = \rho/(1-\rho)$  (Weyl and Fabinger, 2013). Consumer loss is scaled by the keep probability  $1-r$ , and the producer-side denominator involves the effective-revenue pass-through  $\mu_p \rho_c$  rather than the list-price pass-through alone. Production cost, handling cost, and a seller-remitted tax all generate the same incidence split because none of them changes the price at which consumers decide whether to return the product; they shift equilibrium exclusively through the effective marginal revenue  $\mu$ . The buyer-remitted tax, in contrast, alters the refund price to the pre-tax list price, so the return rate and the per-return refund liability differ. This breaks the symmetry, giving an incidence ratio that departs from (25)—a welfare counterpart to the tax-neutrality failure established earlier. When returns are absent ( $r=0$ ,  $\mu_p=1$ ,  $\rho_\tau^b=\rho_c$ ), all four ratios collapse to  $\rho_c/(1-\rho_c)$ , restoring the familiar equivalence.